

# RAG增强智能问答系统 - 详细设计文档

## 1. 引言

### 1.1 文档目的

本文档详细描述RAG增强智能问答系统各模块的内部设计，包括类设计、接口定义、算法详解和数据结构，为编码实现提供直接指导。

### 1.2 参考文档

- 《需求分析.md》
- 《概要设计.md》

## 2. 模块详细设计

### 2.1 配置管理模块 (config.py)

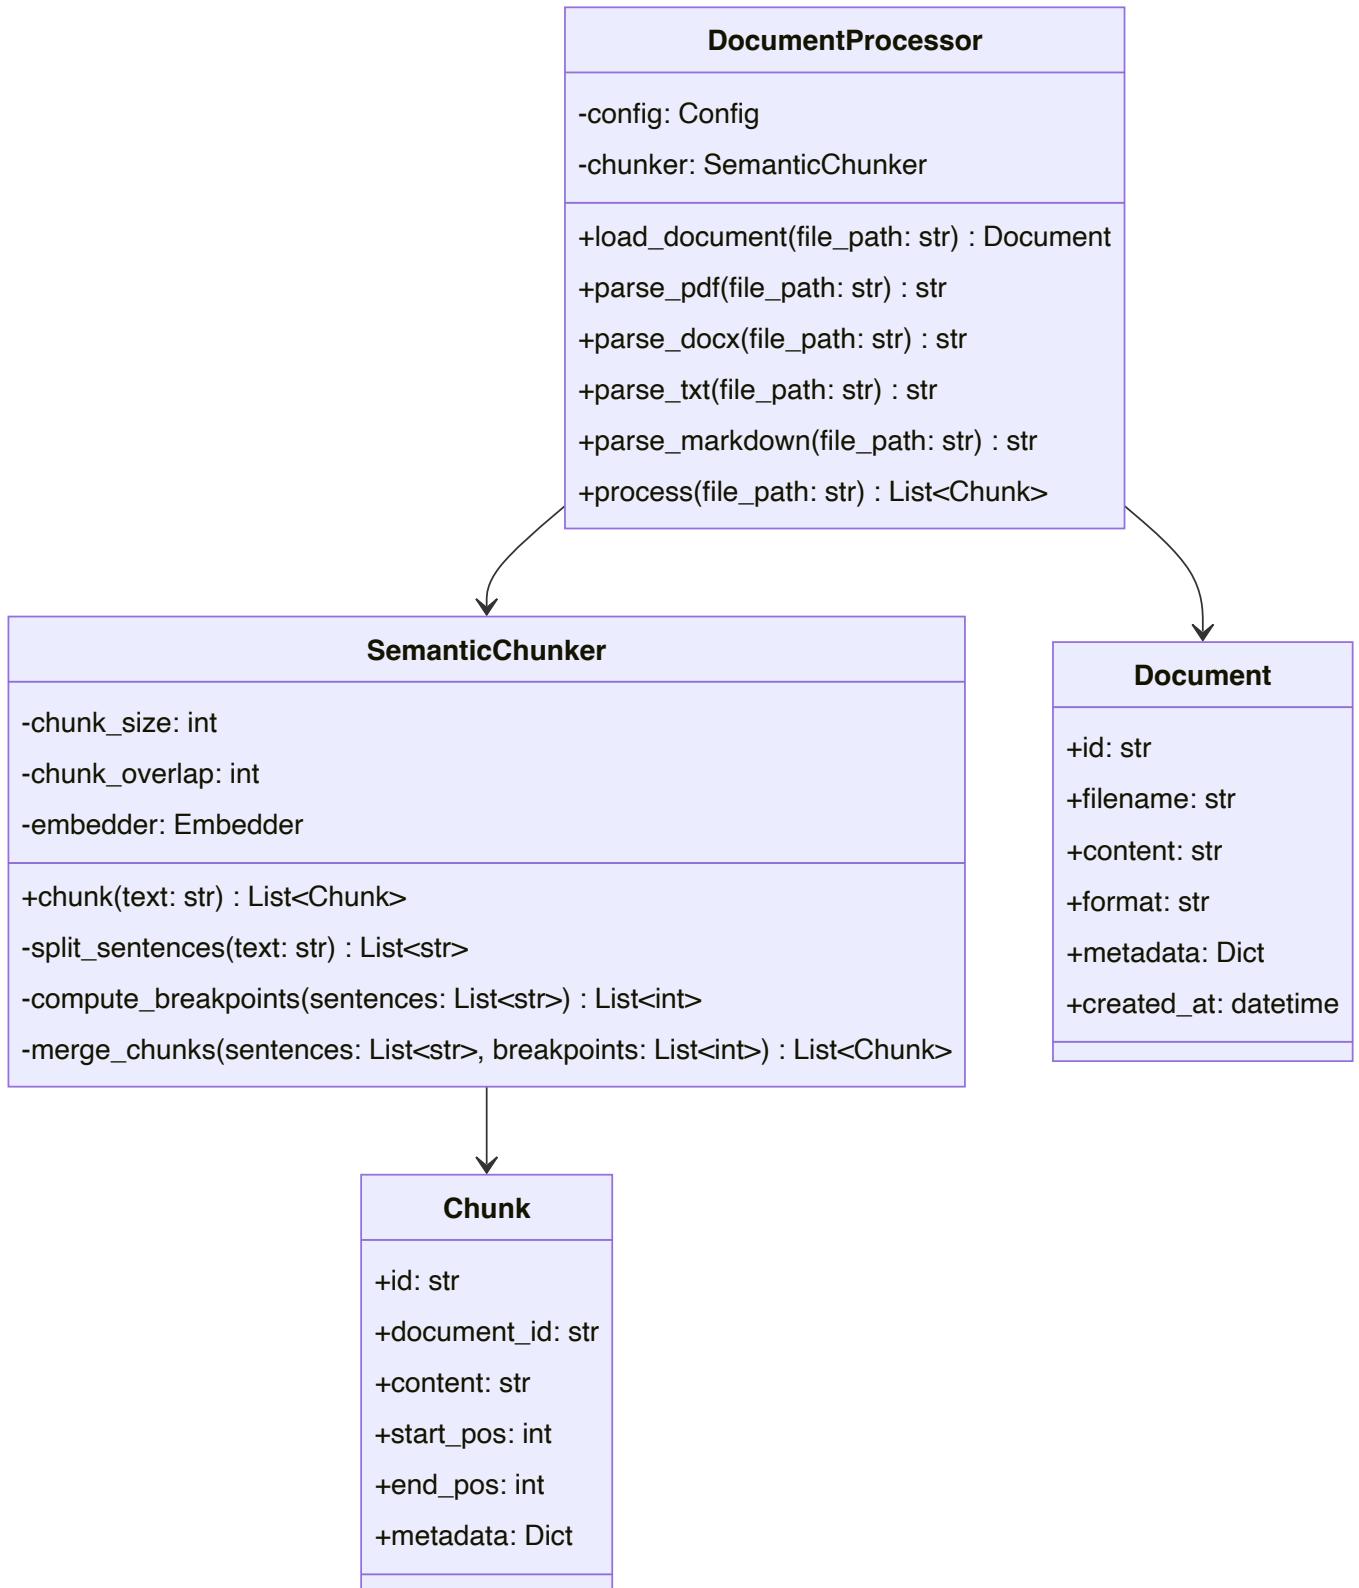
#### 2.1.1 类设计

```
class Config:  
    """系统配置管理类"""\n\n    # 文档处理配置\n    CHUNK_SIZE: int = 512          # 分块大小(tokens)\n    CHUNK_OVERLAP: int = 64         # 分块重叠\n    MAX_FILE_SIZE: int = 50 * 1024 * 1024 # 最大文件大小(50MB)\n    SUPPORTED_FORMATS: List[str] = [".pdf", ".txt", ".docx", ".md"]\n\n    # 嵌入模型配置\n    EMBEDDING_MODEL: str = "BAAI/bge-m3"\n    EMBEDDING_DIM: int = 1024\n    EMBEDDING_BATCH_SIZE: int = 32\n\n    # 向量数据库配置\n    VECTOR_DB_PATH: str = "./data/chroma_db"\n    COLLECTION_NAME: str = "documents"\n\n    # 检索配置\n    TOP_K: int = 5\n    RERANK_TOP_K: int = 10\n    HYBRID_ALPHA: float = 0.7        # 混合检索权重\n\n    # 生成配置\n    LLM_MODEL: str = "qwen2.5:7b"\n    MAX_NEW_TOKENS: int = 1024\n    TEMPERATURE: float = 0.7
```

```
# 服务配置
HOST: str = "0.0.0.0"
PORT: int = 7860
```

## 2.2 文档处理模块 (document\_processor.py)

### 2.2.1 类图



## 2.2.2 核心方法详解

### load\_document()

```

def load_document(self, file_path: str) -> Document:
    """
    加载并解析文档
  
```

```

Args:
    file_path: 文档路径

Returns:
    Document: 解析后的文档对象

Raises:
    ValueError: 不支持的文件格式
    FileNotFoundError: 文件不存在
"""

```

## 处理流程:

1. 验证文件存在性和格式
2. 根据扩展名选择解析器
3. 提取文本内容和元数据
4. 构建Document对象返回

## SemanticChunker.chunk() - 创新算法

```

def chunk(self, text: str, doc_id: str) -> List[Chunk]:
    """
    基于语义边界的智能分块

    算法步骤:
    1. 使用句子分割器切分文本
    2. 计算相邻句子的嵌入向量
    3. 计算相邻句子间的余弦相似度
    4. 找出相似度低于阈值的位置作为潜在分割点
    5. 根据目标块大小合并句子，优先在潜在分割点切分
"""

```

## 算法伪代码:

```

function semantic_chunk(text, target_size, overlap):
    sentences = split_into_sentences(text)
    embeddings = embed_sentences(sentences)

    # 计算相邻句子相似度
    similarities = []
    for i in range(len(sentences) - 1):
        sim = cosine_similarity(embeddings[i], embeddings[i+1])
        similarities.append(sim)

    # 找出语义断点（相似度低的位置）
    threshold = percentile(similarities, 25)
    breakpoints = [i for i, s in enumerate(similarities) if s < threshold]

    # 合并成块
    chunks = []

```

```
current_chunk = []
current_size = 0

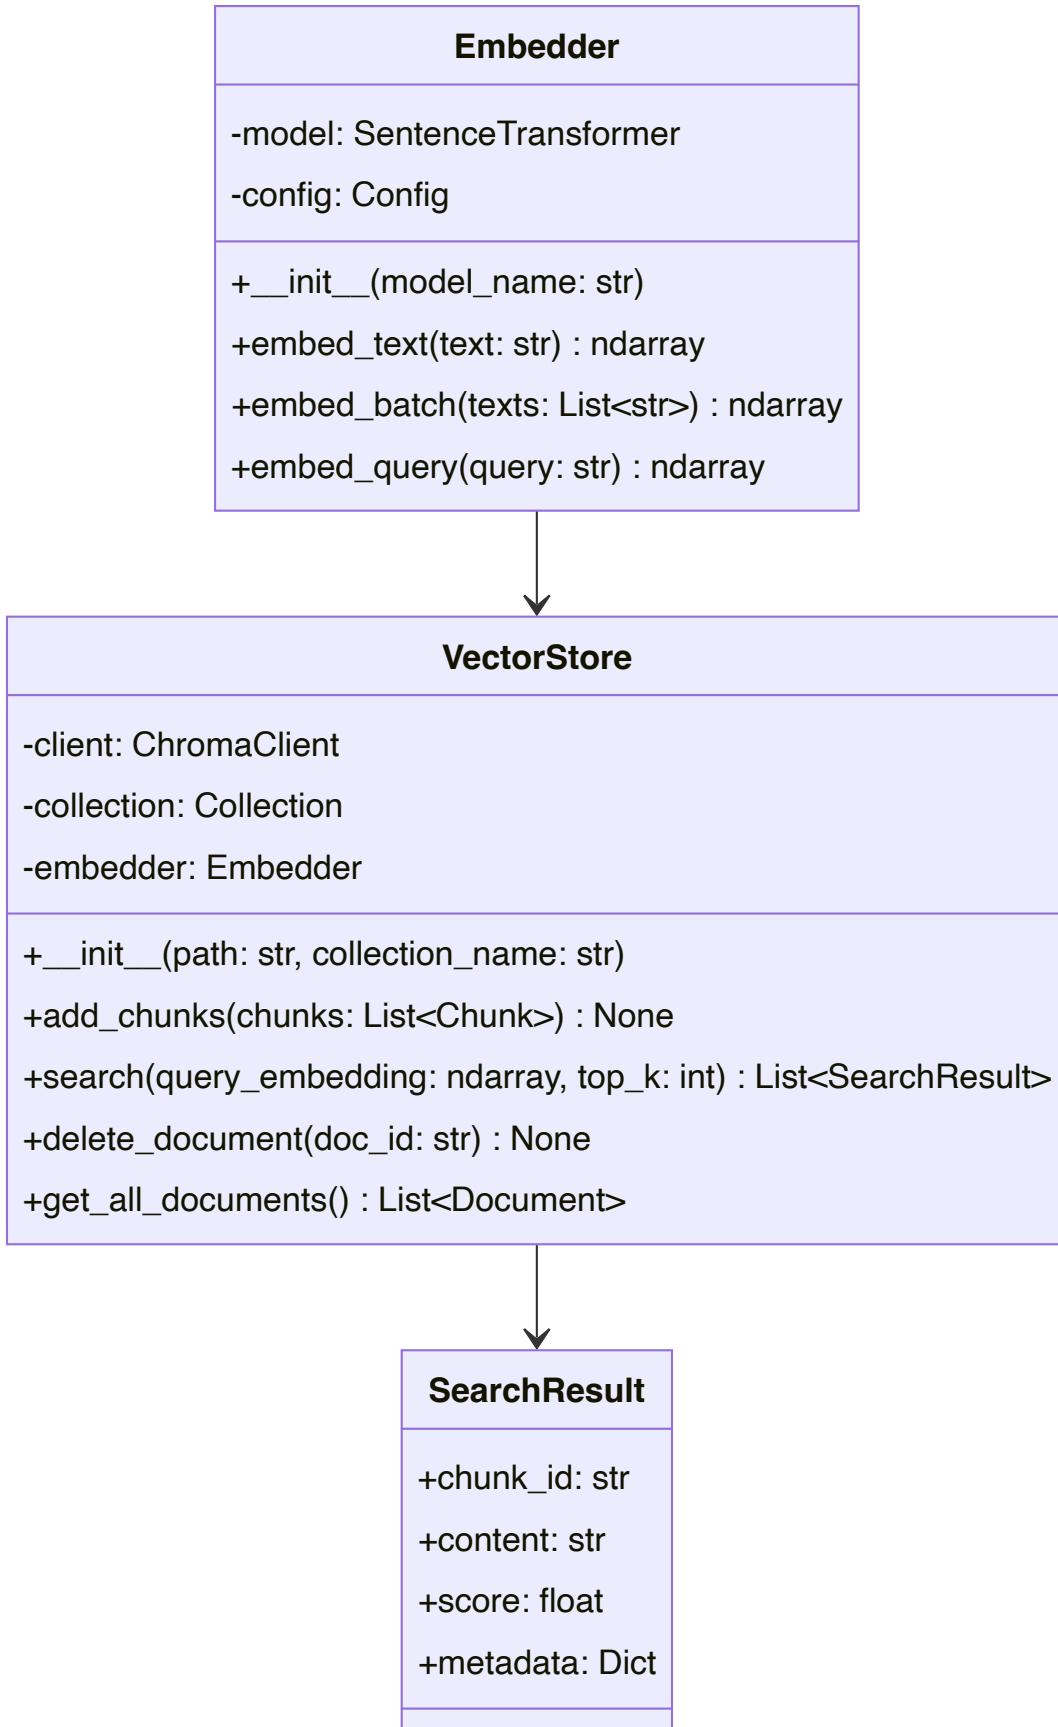
for i, sentence in enumerate(sentences):
    if current_size + len(sentence) > target_size and i in breakpoints:
        chunks.append(join(current_chunk))
        # 保留overlap
        current_chunk = current_chunk[-overlap:]
        current_size = sum(len(s) for s in current_chunk)
    current_chunk.append(sentence)
    current_size += len(sentence)

if current_chunk:
    chunks.append(join(current_chunk))

return chunks
```

## 2.3 嵌入模块 (embedder.py)

### 2.3.1 类图



### 2.3.2 核心方法详解

`embed_batch()`

```

def embed_batch(self, texts: List[str]) -> np.ndarray:
    """
    批量文本嵌入

    Args:
        texts: 文本列表

    Returns:
        embeddings: 形状为 (n, dim) 的嵌入矩阵
    """

   实现细节:
    - 使用模型的encode方法
    - 自动批处理优化内存
    - 归一化向量便于余弦相似度计算
    """

    embeddings = self.model.encode(
        texts,
        batch_size=self.config.EMBEDDING_BATCH_SIZE,
        normalize_embeddings=True,
        show_progress_bar=True
    )
    return embeddings

```

## VectorStore.add\_chunks()

```

def add_chunks(self, chunks: List[Chunk]) -> None:
    """
    添加文本块到向量数据库

    Args:
        chunks: 文本块列表

    实现流程:
    1. 提取文本内容列表
    2. 批量计算嵌入向量
    3. 构建元数据
    4. 调用ChromaDB的add方法
    """

    texts = [chunk.content for chunk in chunks]
    embeddings = self.embedder.embed_batch(texts)

    ids = [chunk.id for chunk in chunks]
    metadatas = [
        {
            "document_id": chunk.document_id,
            "start_pos": chunk.start_pos,
            "end_pos": chunk.end_pos,
            **chunk.metadata
        }
        for chunk in chunks
    ]

```

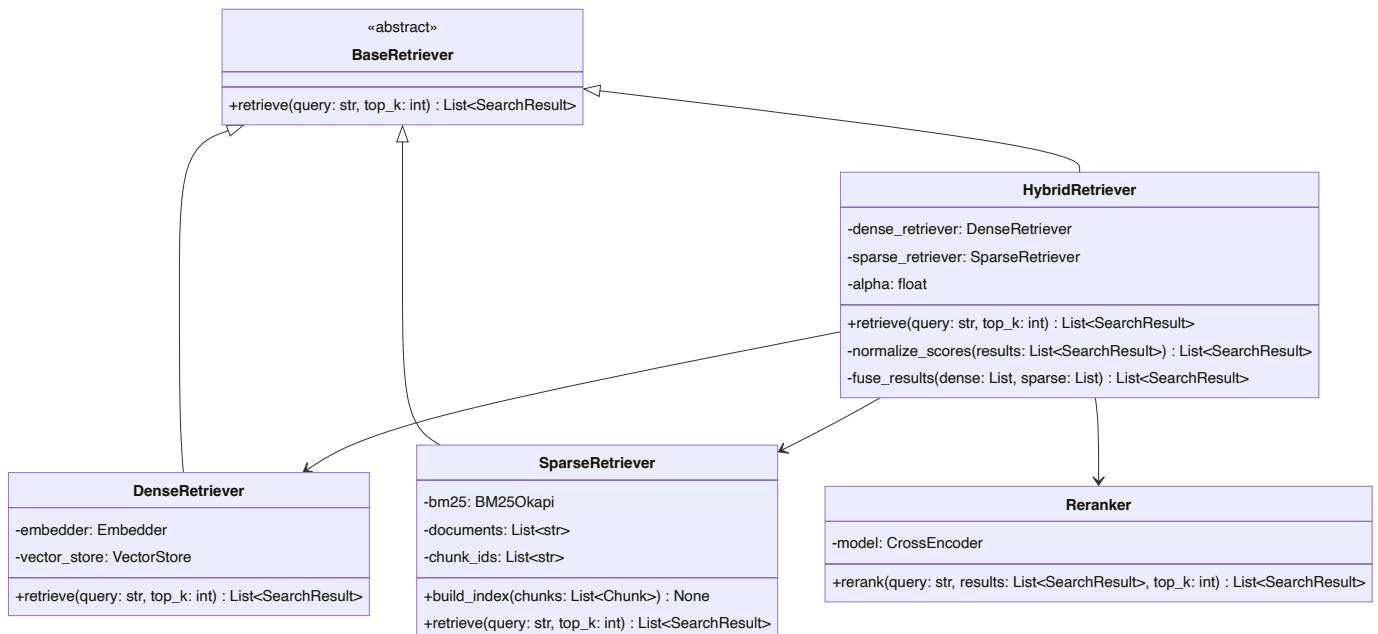
```

self.collection.add(
    ids=ids,
    embeddings=embeddings.tolist(),
    documents=texts,
    metadatas=metadatas
)

```

## 2.4 检索模块 (retriever.py)

### 2.4.1 类图



### 2.4.2 混合检索算法 - 创新点

```
def retrieve(self, query: str, top_k: int = 5) -> List[SearchResult]:
```

```
"""

```

混合检索：结合稠密检索和稀疏检索，带多层过滤

算法：

1. 稠密检索获取  $\text{top\_k} * 2$  结果
2. 预过滤（原始相似度  $\geq$  阈值）
3. 稀疏检索获取  $\text{top\_k} * 2$  结果
4. 分数归一化
5. 加权融合
6. 重排序
7. 最终过滤（融合分数+重排序分数双重验证）

```
"""

```

# 1. 稠密检索

```
dense_results = self.dense_retriever.retrieve(query, top_k * 2)
```

# 2. 预过滤 - 关键！避免归一化后低分结果通过

```
dense_results = self._pre_filter_dense(dense_results)
```

```

# 3. 稀疏检索
sparse_results = self.sparse_retriever.retrieve(query, top_k * 2)

if not dense_results and not sparse_results:
    return [] # 无相关结果, 将触发普通对话模式

# 4. 归一化分数到 [0, 1]
dense_results = self._normalize_scores(dense_results)
sparse_results = self._normalize_scores(sparse_results)

# 5. 加权融合
fused_results = self._fuse_results(dense_results, sparse_results)

# 6. 重排序
if self.reranker and self.reranker.is_available:
    fused_results = self.reranker.rerank(query, fused_results, top_k)

# 7. 最终过滤
fused_results = self._final_filter(fused_results)

return fused_results[:top_k]

def fuse_results(
    self,
    dense: List[SearchResult],
    sparse: List[SearchResult]
) -> List[SearchResult]:
    """
    融合稠密和稀疏检索结果

    公式: final_score = α × dense_score + (1-α) × sparse_score
    """
    score_map = {}
    result_map = {}

    # 收集稠密检索分数
    for r in dense:
        score_map[r.chunk_id] = {"dense": r.score, "sparse": 0}
        result_map[r.chunk_id] = r

    # 收集稀疏检索分数
    for r in sparse:
        if r.chunk_id in score_map:
            score_map[r.chunk_id]["sparse"] = r.score
        else:
            score_map[r.chunk_id] = {"dense": 0, "sparse": r.score}
            result_map[r.chunk_id] = r

    # 计算融合分数
    fused = []

```

```

for chunk_id, scores in score_map.items():
    final_score = (
        self.alpha * scores["dense"] +
        (1 - self.alpha) * scores["sparse"]
    )
    result = result_map[chunk_id]
    result.score = final_score
    fused.append(result)

# 按分数降序排序
fused.sort(key=lambda x: x.score, reverse=True)
return fused

```

### 2.4.3 多层过滤算法 - 创新点

```

def _pre_filter_dense(self, results: List[SearchResult]) -> List[SearchResult]:
    """
    预过滤：在归一化前用原始余弦相似度过滤

    目的：避免低相关结果因归一化而通过阈值
    阈值：0.5（余弦相似度范围 [-1, 1]）
    """
    return [r for r in results if r.score >= self.similarity_threshold]

def _final_filter(self, results: List[SearchResult]) -> List[SearchResult]:
    """
    最终过滤：同时检查融合分数和重排序分数

    策略：
    - 有重排序分数时：rerank_score > 0 且 score >= 0.25
    - 无重排序分数时：score >= 0.25
    """
    filtered = []
    for r in results:
        if hasattr(r, 'rerank_score') and r.rerank_score is not None:
            # 两个条件都要满足
            if r.rerank_score > 0 and r.score >= 0.25:
                filtered.append(r)
        else:
            if r.score >= 0.25:
                filtered.append(r)
    return filtered

```

#### 设计原理：

- 问“你好”等无关问题时，虽然向量检索可能返回结果，但原始相似度很低 (< 0.5)
- 预过滤在归一化前执行，确保低相关结果不会因归一化而变成高分
- 最终过滤双重验证，确保返回的结果真正相关

### 2.4.4 重排序算法 - 创新点

```

class Reranker:
    """使用Cross-Encoder进行重排序"""

    def __init__(self, model_name: str = "BAAI/bge-reranker-base"):
        self.model = CrossEncoder(model_name)

    def rerank(
        self,
        query: str,
        results: List[SearchResult],
        top_k: int
    ) -> List[SearchResult]:
        """
        对检索结果重排序
        """

        原理:
        Cross-Encoder同时接收query和document作为输入,
        能够建模更精细的交互关系, 排序效果优于Bi-Encoder
        """

        if not results:
            return results

        # 构建输入对
        pairs = [(query, r.content) for r in results]

        # 获取重排序分数
        scores = self.model.predict(pairs)

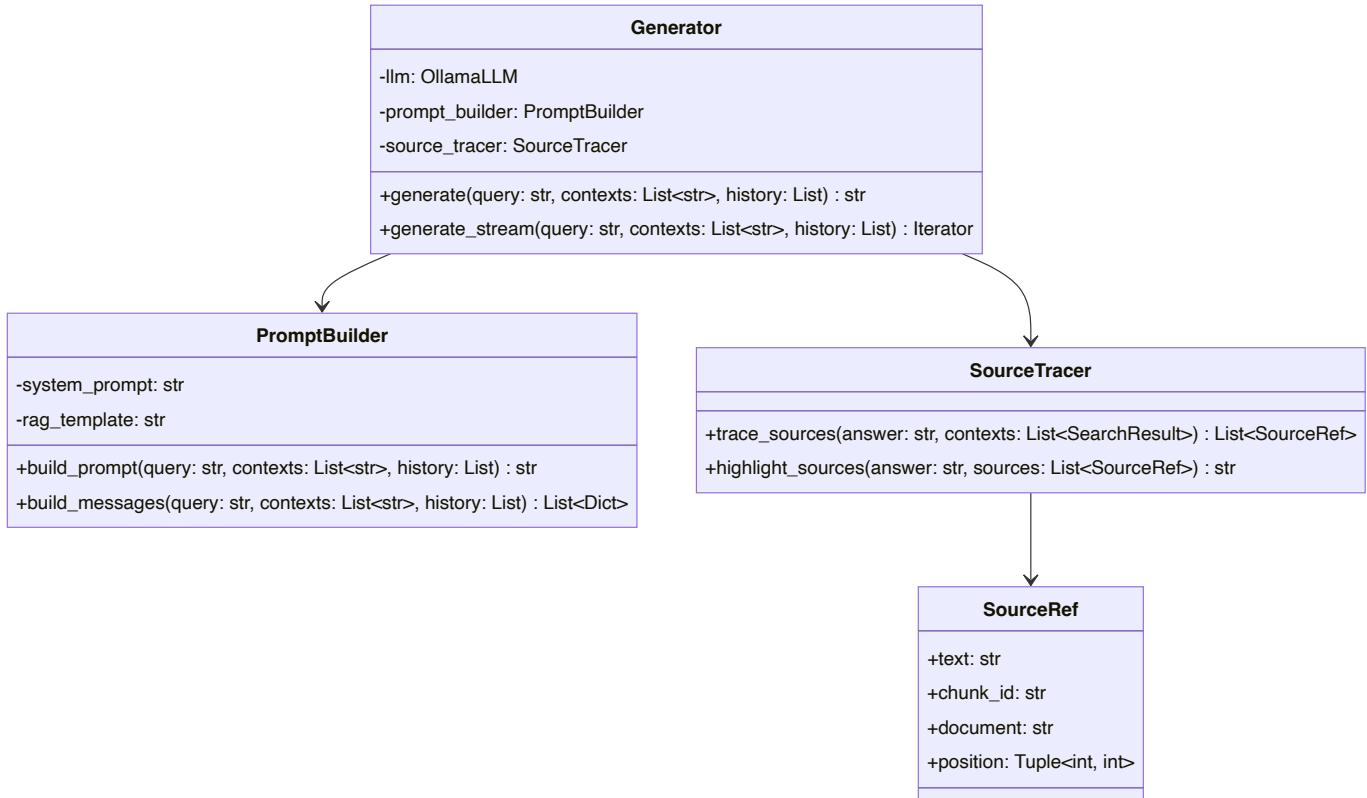
        # 更新分数并排序
        for i, result in enumerate(results):
            result.rerank_score = float(scores[i])

        results.sort(key=lambda x: x.rerank_score, reverse=True)
        return results[:top_k]

```

## 2.5 生成模块 (generator.py)

### 2.5.1 类图



## 2.5.2 Prompt模板设计

系统根据是否有检索结果自动切换Prompt模式：

```
# RAG模式系统提示（有知识库时）
SYSTEM_PROMPT_RAG = """你是一个专业的问答助手。请根据提供的参考资料回答用户问题。
```

规则：

1. 优先根据参考资料回答，不要编造信息
2. 如果参考资料中没有相关内容，请明确说明“根据现有资料无法回答”
3. 回答时引用来源，使用[1], [2]等标记
4. 回答要准确、简洁、有条理
5. 使用与用户问题相同的语言回答

"""

```
# 普通对话模式系统提示（无知识库时）
```

```
SYSTEM_PROMPT_CHAT = """你是一个友好、专业的AI助手。请用准确、简洁、有条理的方式回答用户问题。
```

规则：

1. 回答要准确、简洁、有条理
2. 使用与用户问题相同的语言回答
3. 如果不确定答案，请诚实说明
4. 保持友好和专业的态度

"""

```
# 动态选择提示词
```

```
def build_messages(self, query, search_results, history):
    if search_results:
        system_prompt = SYSTEM_PROMPT_RAG
```

```

user_message = f"参考资料: {contexts}\n\n问题: {query}"
else:
    system_prompt = SYSTEM_PROMPT_CHAT # 无知识库, 普通对话
    user_message = query
...

```

### 2.5.3 LLM提供商抽象 - 创新点

```

class Generator:
    """支持多种LLM后端的答案生成器"""

    def __init__(
        self,
        provider: str = "ollama", # ollama 或 openai
        model: str = "qwen2.5:7b",
        base_url: str = "http://localhost:11434",
        api_key: str = None, # OpenAI模式需要
        ...
    ):
        self.provider = provider
        self._init_client()

    def _init_client(self):
        if self.provider == "ollama":
            self._init_ollama()
        elif self.provider == "openai":
            self._init_openai()

    def _init_openai(self):
        """支持OpenAI兼容API (DeepSeek、Azure等) """
        from openai import OpenAI
        client_kwargs = {"api_key": self.api_key}
        if self.base_url: # 自定义API地址
            client_kwargs["base_url"] = self.base_url
        self._client = OpenAI(**client_kwargs)

```

### 2.5.3 来源追踪算法 - 创新点

```

class SourceTracer:
    """答案来源追踪器"""

    def trace_sources(
        self,
        answer: str,
        contexts: List[SearchResult]
    ) -> List[SourceRef]:
        """
        追踪答案中的内容来源
        """

        算法:
        1. 提取答案中的关键句子

```

```

2. 与每个context计算相似度
3. 相似度超过阈值则标记为来源
"""

sources = []
answer_sentences = self._split_sentences(answer)

for sent in answer_sentences:
    best_match = None
    best_score = 0

    for ctx in contexts:
        # 计算句子与context的相似度
        score = self._compute_similarity(sent, ctx.content)
        if score > best_score and score > 0.6:
            best_score = score
            best_match = ctx

    if best_match:
        sources.append(SourceRef(
            text=sent,
            chunk_id=best_match.chunk_id,
            document=best_match.metadata.get("filename", "unknown"),
            score=best_score
        ))

```

return self.\_deduplicate(sources)

```

def highlight_sources(
    self,
    answer: str,
    sources: List[SourceRef]
) -> str:
"""

```

在答案中添加来源标记

输出格式:

```

"机器学习是一种人工智能方法[1]..."
"""

# 按出现位置排序来源
source_map = {}
for i, source in enumerate(sources, 1):
    if source.chunk_id not in source_map:
        source_map[source.chunk_id] = i

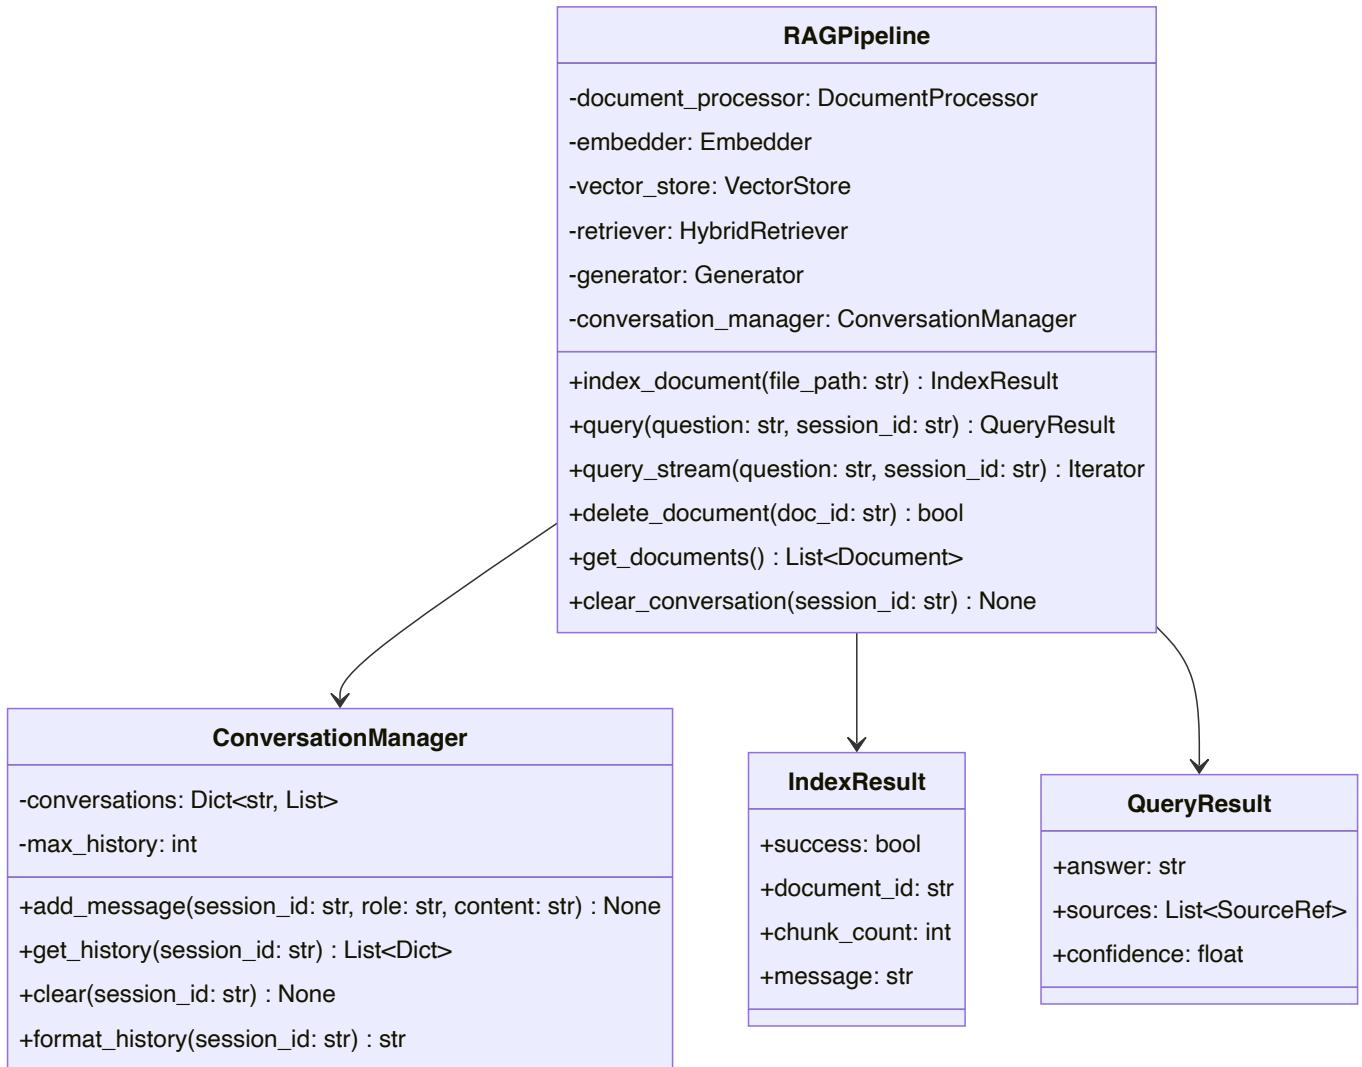
# 添加引用标记
highlighted = answer
for source in sources:
    ref_num = source_map[source.chunk_id]
    # 在对应句子后添加引用
    highlighted = highlighted.replace(
        source.text,
        f"{source.text}[{ref_num}]"

```

```
    )
return highlighted
```

## 2.6 RAG流水线 (rag\_pipeline.py)

### 2.6.1 类图



### 2.6.2 核心流程

#### 文档索引流程

```
def index_document(self, file_path: str) -> IndexResult:
```

```
    """
```

索引文档

流程：

1. 验证文件
2. 加载并解析文档
3. 智能分块

```

4. 向量化并存储
5. 更新稀疏索引
"""

try:
    # 1. 加载文档
    document = self.document_processor.load_document(file_path)

    # 2. 智能分块
    chunks = self.document_processor.process(document)

    # 3. 向量化存储
    self.vector_store.add_chunks(chunks)

    # 4. 更新BM25索引
    self.retriever.sparse_retriever.add_chunks(chunks)

    return IndexResult(
        success=True,
        document_id=document.id,
        chunk_count=len(chunks),
        message=f"成功索引文档, 共{len(chunks)}个文本块"
    )
except Exception as e:
    return IndexResult(
        success=False,
        document_id=None,
        chunk_count=0,
        message=f"索引失败: {str(e)}"
)

```

## 问答流程

```

def query(self, question: str, session_id: str = "default") -> QueryResult:
    """
    处理用户问题

    流程:
    1. 获取对话历史
    2. 问题改写 (可选, 结合历史理解指代)
    3. 混合检索
    4. 重排序
    5. 生成答案
    6. 来源追踪
    7. 更新对话历史
    """

    # 1. 获取对话历史
    history = self.conversation_manager.get_history(session_id)

    # 2. 检索相关内容
    search_results = self.retriever.retrieve(
        question,
        top_k=self.config.TOP_K

```

```

)
if not search_results:
    return QueryResult(
        answer="抱歉，根据现有知识库无法找到相关内容。",
        sources=[],
        confidence=0.0
)

# 3. 构建上下文
contexts = [r.content for r in search_results]

# 4. 生成答案
answer = self.generator.generate(
    query=question,
    contexts=contexts,
    history=history
)

# 5. 来源追踪
sources = self.generator.source_tracer.trace_sources(
    answer,
    search_results
)

# 6. 添加引用标记
answer_with_refs = self.generator.source_tracer.highlight_sources(
    answer,
    sources
)

# 7. 更新对话历史
self.conversation_manager.add_message(session_id, "user", question)
self.conversation_manager.add_message(session_id, "assistant", answer)

# 8. 计算置信度
confidence = sum(r.score for r in search_results[:3]) / 3

return QueryResult(
    answer=answer_with_refs,
    sources=sources,
    confidence=confidence
)

```

## 2.7 API服务 (api.py)

### 2.7.1 接口详细定义

```

from fastapi import FastAPI, UploadFile, File, HTTPException
from pydantic import BaseModel

```

```

from typing import List, Optional

app = FastAPI(title="RAG问答系统API")

# ===== 数据模型 =====

class QueryRequest(BaseModel):
    question: str
    session_id: Optional[str] = "default"
    top_k: Optional[int] = 5

class QueryResponse(BaseModel):
    answer: str
    sources: List[dict]
    confidence: float

class DocumentInfo(BaseModel):
    id: str
    filename: str
    chunk_count: int
    created_at: str

class UploadResponse(BaseModel):
    status: str
    document_id: str
    chunk_count: int
    message: str

# ===== API端点 =====

@app.post("/api/documents/upload", response_model=UploadResponse)
async def upload_document(file: UploadFile = File(...)):
    """
    上传并索引文档

    - 支持格式: PDF, TXT, DOCX, Markdown
    - 最大文件大小: 50MB
    """

    # 验证文件格式
    if not any(file.filename.endswith(ext) for ext in SUPPORTED_FORMATS):
        raise HTTPException(400, "不支持的文件格式")

    # 保存文件
    file_path = save_upload_file(file)

    # 索引文档
    result = rag_pipeline.index_document(file_path)

    if not result.success:
        raise HTTPException(500, result.message)

    return UploadResponse(

```

```

        status="success",
        document_id=result.document_id,
        chunk_count=result.chunk_count,
        message=result.message
    )

@app.get("/api/documents", response_model=List[DocumentInfo])
async def list_documents():
    """获取已索引的文档列表"""
    documents = rag_pipeline.get_documents()
    return [
        DocumentInfo(
            id=doc.id,
            filename=doc.filename,
            chunk_count=doc.chunk_count,
            created_at=doc.created_at.isoformat()
        )
        for doc in documents
    ]

@app.delete("/api/documents/{doc_id}")
async def delete_document(doc_id: str):
    """删除指定文档"""
    success = rag_pipeline.delete_document(doc_id)
    if not success:
        raise HTTPException(404, "文档不存在")
    return {"status": "success", "message": "文档已删除"}

@app.post("/api/qa/query", response_model=QueryResponse)
async def query(request: QueryRequest):
    """
    提交问题并获取答案
    """

    - 支持多轮对话（通过session_id关联）
    - 返回答案及来源引用
    """

    result = rag_pipeline.query(
        question=request.question,
        session_id=request.session_id
    )

    return QueryResponse(
        answer=result.answer,
        sources=[
            {
                "chunk_id": s.chunk_id,
                "document": s.document,
                "text": s.text,
                "score": s.score
            }
            for s in result.sources
        ],
    )

```

```

confidence=result.confidence
)

@app.post("/api/conversation/clear")
async def clear_conversation(session_id: str = "default"):
    """清空对话历史"""
    rag_pipeline.clear_conversation(session_id)
    return {"status": "success", "message": "对话历史已清空"}

```

## 2.8 前端应用 (gradio\_app.py)

### 2.8.1 界面设计



### 2.8.2 组件设计

```

import gradio as gr

def create_app():
    with gr.Blocks(title="RAG智能问答系统", theme=gr.themes.Soft()) as app:
        gr.Markdown("# 🤖 RAG增强智能问答系统")
        gr.Markdown("上传文档，然后基于文档内容进行问答")

        with gr.Row():
            # 左侧面板
            with gr.Column(scale=1):
                gr.Markdown("### 📄 文档管理")

                file_upload = gr.File(
                    label="上传文档",
                    file_types=[".pdf", ".txt", ".docx", ".md"],
                    file_count="multiple"

```

```

)
upload_btn = gr.Button("📤 上传并索引", variant="primary")
upload_status = gr.Textbox(label="上传状态", interactive=False)

gr.Markdown("### 📄 已索引文档")
doc_list = gr.Dataframe(
    headers=["文档名", "块数", "操作"],
    label="文档列表"
)
refresh_btn = gr.Button("⟳ 刷新列表")

gr.Markdown("### 🛡️ 设置")
top_k_slider = gr.Slider(1, 10, value=5, step=1, label="检索数量")
clear_btn = gr.Button("🧹 清空对话")

# 右侧面板
with gr.Column(scale=2):
    gr.Markdown("### 💬 问答对话")

    chatbot = gr.Chatbot(
        label="对话历史",
        height=400,
        show_label=False
    )

    with gr.Row():
        question_input = gr.Textbox(
            label="输入问题",
            placeholder="请输入您的问题...",
            scale=4
        )
        send_btn = gr.Button("发送", variant="primary", scale=1)

    gr.Markdown("### 📚 来源引用")
    sources_display = gr.JSON(label="答案来源")

# 事件绑定
upload_btn.click(
    fn=handle_upload,
    inputs=[file_upload],
    outputs=[upload_status, doc_list]
)

send_btn.click(
    fn=handle_query,
    inputs=[question_input, chatbot, top_k_slider],
    outputs=[chatbot, sources_display, question_input]
)

question_input.submit(
    fn=handle_query,
    inputs=[question_input, chatbot, top_k_slider],
)

```

```

        outputs=[chatbot, sources_display, question_input]
    )

    clear_btn.click(
        fn=handle_clear,
        inputs=[],
        outputs=[chatbot, sources_display]
    )

    refresh_btn.click(
        fn=handle_refresh,
        inputs=[],
        outputs=[doc_list]
    )

return app

```

## 3. 数据结构详细设计

### 3.1 核心数据类

```

from dataclasses import dataclass, field
from typing import Dict, List, Optional, Tuple
from datetime import datetime
import uuid

@dataclass
class Document:
    """文档数据类"""
    id: str = field(default_factory=lambda: str(uuid.uuid4()))
    filename: str = ""
    content: str = ""
    format: str = ""
    metadata: Dict = field(default_factory=dict)
    created_at: datetime = field(default_factory=datetime.now)
    chunk_count: int = 0

@dataclass
class Chunk:
    """文本块数据类"""
    id: str = field(default_factory=lambda: str(uuid.uuid4()))
    document_id: str = ""
    content: str = ""
    start_pos: int = 0
    end_pos: int = 0
    metadata: Dict = field(default_factory=dict)

@dataclass
class SearchResult:
    """检索结果数据类"""

```

```

chunk_id: str
content: str
score: float
metadata: Dict
rerank_score: Optional[float] = None

@dataclass
class SourceRef:
    """来源引用数据类"""
    text: str
    chunk_id: str
    document: str
    score: float
    position: Optional[Tuple[int, int]] = None

@dataclass
class QueryResult:
    """查询结果数据类"""
    answer: str
    sources: List[SourceRef]
    confidence: float

@dataclass
class IndexResult:
    """索引结果数据类"""
    success: bool
    document_id: Optional[str]
    chunk_count: int
    message: str

```

## 4. 错误处理设计

### 4.1 自定义异常

```

class RAGEException(Exception):
    """RAG系统基础异常"""
    pass

class DocumentParseError(RAGEException):
    """文档解析错误"""
    pass

class EmbeddingError(RAGEException):
    """嵌入计算错误"""
    pass

class RetrievalError(RAGEException):
    """检索错误"""
    pass

```

```

class GenerationError(RAGEException):
    """生成错误"""
    pass

class ConfigError(RAGEException):
    """配置错误"""
    pass

```

## 4.2 错误处理策略

```

def safe_execute(func, *args, fallback=None, **kwargs):
    """安全执行函数，捕获异常并返回fallback"""
    try:
        return func(*args, **kwargs)
    except RAGEException as e:
        logger.error(f"RAG错误: {e}")
        return fallback
    except Exception as e:
        logger.exception(f"未知错误: {e}")
        return fallback

```

## 5. 日志设计

```

import logging

def setup_logging():
    """配置日志系统"""
    logging.basicConfig(
        level=logging.INFO,
        format='%(asctime)s - %(name)s - %(levelname)s - %(message)s',
        handlers=[
            logging.FileHandler('rag_system.log'),
            logging.StreamHandler()
        ]
    )

    # 设置各模块日志级别
    logging.getLogger('document_processor').setLevel(logging.INFO)
    logging.getLogger('retriever').setLevel(logging.INFO)
    logging.getLogger('generator').setLevel(logging.INFO)

```

## 附录A: 配置文件示例

```

# config.yaml
document:
    chunk_size: 512
    chunk_overlap: 64

```

```

max_file_size: 52428800 # 50MB
supported_formats:
  - .pdf
  - .txt
  - .docx
  - .md

embedding:
  model: "BAAI/bge-m3"
  dimension: 1024
  batch_size: 32

retrieval:
  top_k: 5
  rerank_top_k: 10
  hybrid_alpha: 0.7

generation:
  model: "qwen2.5:7b"
  max_tokens: 1024
  temperature: 0.7

server:
  host: "0.0.0.0"
  port: 7860

```

## 附录B: 依赖列表

```

# requirements.txt
langchain>=0.1.0
langchain-community>=0.0.10
chromadb>=0.4.0
sentence-transformers>=2.2.0
transformers>=4.35.0
torch>=2.0.0
gradio>=4.0.0
fastapi>=0.100.0
uvicorn>=0.23.0
python-multipart>=0.0.6
pypdf>=3.0.0
python-docx>=0.8.11
markdown>=3.4.0
rank-bm25>=0.2.2
numpy>=1.24.0
pydantic>=2.0.0
ollama>=0.1.0

```